EXP NO: 10	Image Generation Using Generative Adversarial Network (Gan)

AIM:

To generate handwritten-digit images using a Generative Adversarial Network (GAN) trained on the MNIST dataset.

ALGORITHM:

1. Load the MNIST training images, scale them to the range -1 to 1, shuffle them and create batches of 128. Build
2. the Generator: Dense(7x7x256) with BatchNormalization and LeakyReLU, reshape, then Conv2DTranspose layers (128, 64, 1 filters) with a tanh output that produces a 28 x 28 x 1 image from a 100-dimensional noise vector.
3. Build the Discriminator: two Conv2D layers (64, 128 filters) with LeakyReLU and Dropout(0.3), followed by Flatten and a Dense(1) output.
4. Define the binary cross-entropy based generator loss and discriminator loss (real loss + fake loss) and an Adam optimizer (learning rate 0.0002, beta_1 = 0.5) for each network.
5. In every training step generate fake images from random noise, compute the discriminator outputs for real and fake images, calculate both losses with GradientTape and update both networks.
6. Train for 10 epochs, print the average generator and discriminator loss, and display images generated from a fixed noise seed at epoch 1 and every second epoch.

CODE:

```
import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt
import os

print("TensorFlow version:", tf.__version__)

(x_train, _), (_, _) = tf.keras.datasets.mnist.load_data()

x_train = x_train.astype("float32")
x_train = (x_train - 127.5) / 127.5
x_train = np.expand_dims(x_train, axis=-1)

BUFFER_SIZE = 60000
BATCH_SIZE = 128
EPOCHS = 10
NOISE_DIM = 100

dataset = tf.data.Dataset.from_tensor_slices(x_train)
dataset = dataset.shuffle(BUFFER_SIZE).batch(BATCH_SIZE, drop_remainder=True)

def build_generator():

    model = tf.keras.Sequential([
        tf.keras.layers.Dense(7 * 7 * 256, use_bias=False, input_shape=(NOISE_DIM,)),
        tf.keras.layers.BatchNormalization(),
        tf.keras.layers.LeakyReLU(),

        tf.keras.layers.Reshape((7, 7, 256)),

        tf.keras.layers.Conv2DTranspose(
        128, (5, 5), strides=(1, 1),
        padding="same", use_bias=False
```

```
            ),
            tf.keras.layers.BatchNormalization(),
            tf.keras.layers.LeakyReLU(),

            tf.keras.layers.Conv2DTranspose(
                  64, (5, 5), strides=(2, 2),
                  padding="same", use_bias=False
            ),
            tf.keras.layers.BatchNormalization(),
            tf.keras.layers.LeakyReLU(),

            tf.keras.layers.Conv2DTranspose(
                  1, (5, 5), strides=(2, 2),
                  padding="same", use_bias=False,
                  activation="tanh"
            )
      ])

      return model

def build_discriminator():
      model = tf.keras.Sequential([
            tf.keras.layers.Conv2D(
                  64, (5, 5), strides=(2, 2),
                  padding="same",
                  input_shape=[28, 28, 1]
            ),
            tf.keras.layers.LeakyReLU(),
            tf.keras.layers.Dropout(0.3),

            tf.keras.layers.Conv2D(
                  128, (5, 5), strides=(2, 2),
                  padding="same"
            ),
            tf.keras.layers.LeakyReLU(),
            tf.keras.layers.Dropout(0.3),

            tf.keras.layers.Flatten(),
            tf.keras.layers.Dense(1)
      ])

      return model

generator = build_generator()
discriminator = build_discriminator()


cross_entropy = tf.keras.losses.BinaryCrossentropy(from_logits=True)

def generator_loss(fake_output):
      return cross_entropy(
            tf.ones_like(fake_output),
            fake_output
      )

def discriminator_loss(real_output, fake_output):
      real_loss = cross_entropy(
            tf.ones_like(real_output),
            real_output
      )

      fake_loss = cross_entropy(
            tf.zeros_like(fake_output),
            fake_output
```

```
    )

    return real_loss + fake_loss

generator_optimizer = tf.keras.optimizers.Adam(
    learning_rate=0.0002,
    beta_1=0.5
)

discriminator_optimizer = tf.keras.optimizers.Adam(
    learning_rate=0.0002,
    beta_1=0.5
)

seed = tf.random.normal([16, NOISE_DIM])
@tf.function

def train_step(images):
    noise = tf.random.normal(
        [BATCH_SIZE, NOISE_DIM]
    )

    withtf.GradientTape() as gen_tape, tf.GradientTape() as disc_tape:

        generated_images = generator(
            noise,
            training=True
        )

        real_output = discriminator(
            images,
            training=True
        )

        fake_output = discriminator(
            generated_images,
            training=True
        )

        gen_loss = generator_loss(fake_output)

        disc_loss = discriminator_loss(
            real_output,
            fake_output
        )

    gradients_of_generator = gen_tape.gradient(
        gen_loss,
        generator.trainable_variables
    )

    gradients_of_discriminator = disc_tape.gradient(
        disc_loss,
        discriminator.trainable_variables
    )

    generator_optimizer.apply_gradients(
        zip(
            gradients_of_generator,
            generator.trainable_variables
        )
    )
```

```
        discriminator_optimizer.apply_gradients(
            zip(
                gradients_of_discriminator,
                discriminator.trainable_variables
            )
        )

        return gen_loss, disc_loss

def generate_images(model, epoch, test_input):
        predictions = model(
            test_input,
            training=False
        )

        fig = plt.figure(figsize=(8, 8))

        for i in range(predictions.shape[0]):
            plt.subplot(4, 4, i + 1)

            plt.imshow(
                predictions[i, :, :, 0] * 0.5 + 0.5,
                cmap="gray"
            )

            plt.axis("off")

        plt.suptitle(
            "Generated Images - Epoch " + str(epoch)
        )

        plt.show()

def train(dataset, epochs):
        for epoch in range(1, epochs + 1):

            total_gen_loss = 0
            total_disc_loss = 0
            batches = 0

            for image_batch in dataset:

                gen_loss, disc_loss = train_step(
                    image_batch
                )

                total_gen_loss += gen_loss
                total_disc_loss += disc_loss
                batches += 1

            avg_gen_loss = total_gen_loss / batches
            avg_disc_loss = total_disc_loss / batches

            print(
                "Epoch {}/{} | Generator Loss: {:.4f} | Discriminator Loss: {:.4f}".format(
                    epoch,  epochs,
                    avg_gen_loss,
                    avg_disc_loss

                )
            )

            if epoch == 1 or epoch % 2 == 0:
```

```
                generate_images(
                    generator,
                    epoch,
                    seed
                )

print("Starting GAN training...\n")

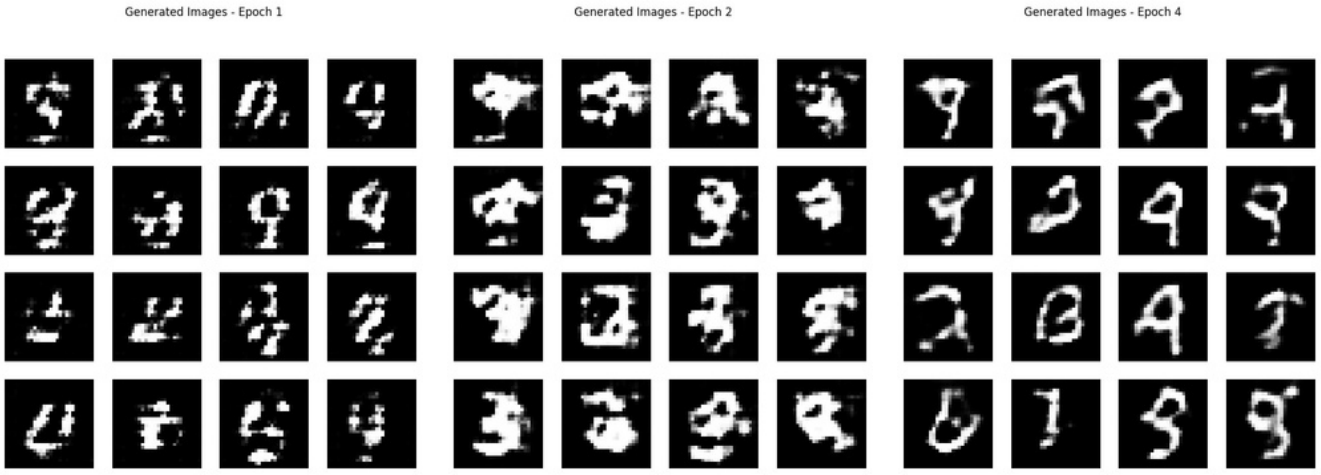
train(
    dataset,
    EPOCHS
)

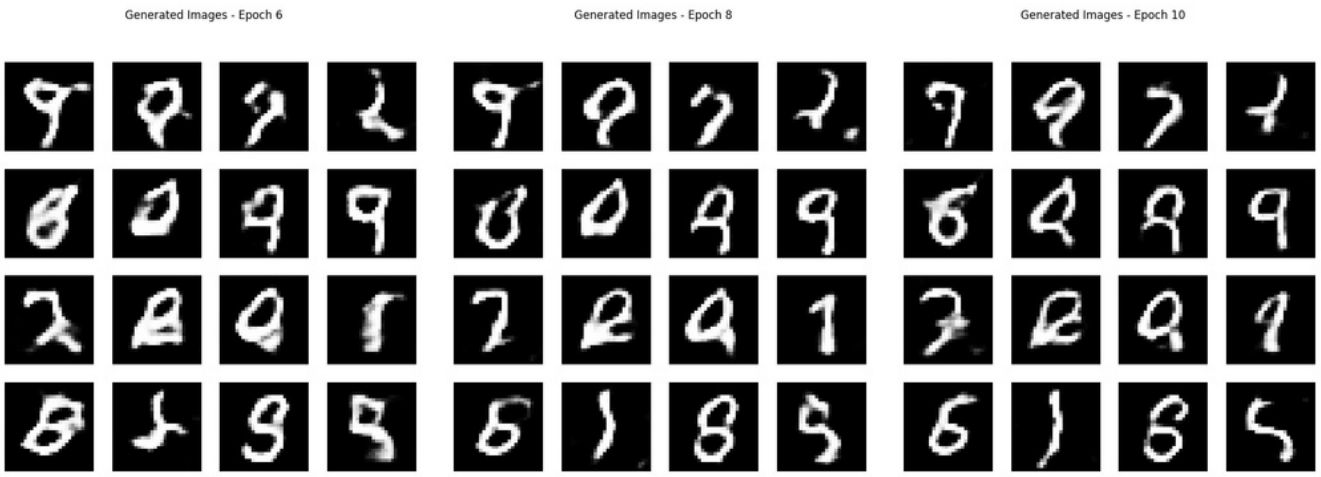
print("\nFinal Generated Images")
generate_images(

    generator,
    EPOCHS,
    seed
)
```

OUTPUT:

```
TensorFlow version: 2.21.0
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz
11490434/11490434 ━━━━━━━━━━━━━━━━━━━━ 0s 0us/step
Starting GAN training...
Epoch 1/10 | Generator Loss: 0.7075 | Discriminator Loss: 1.3177
Epoch 2/10 | Generator Loss: 0.7592 | Discriminator Loss: 1.3262
Epoch 3/10 | Generator Loss: 0.8213 | Discriminator Loss: 1.2745
Epoch 4/10 | Generator Loss: 0.8356 | Discriminator Loss: 1.2669
Epoch 5/10 | Generator Loss: 0.8073 | Discriminator Loss: 1.3021
Epoch 6/10 | Generator Loss: 0.7930 | Discriminator Loss: 1.3155
Epoch 7/10 | Generator Loss: 0.7892 | Discriminator Loss: 1.3183
Epoch 8/10 | Generator Loss: 0.7854 | Discriminator Loss: 1.3161
Epoch 9/10 | Generator Loss: 0.8037 | Discriminator Loss: 1.3062
Epoch 10/10 | Generator Loss: 0.7957 | Discriminator Loss: 1.3092
```





RESULT:

The GAN was trained on MNIST for 10 epochs, ending with a generator loss of 0.7957 and a discriminator loss of 1.3092. The generated images improved from noise-like patterns in epoch 1 to digit-like shapes by epoch 10.